

Operating Systems Internals – Task scheduling 3

Péter Györke


<http://www.mit.bme.hu/~gyorke/>

gyorke@mit.bme.hu

Budapest University of Technology and Economics (BME)
Department of Measurement and Information Systems (MIT)

The slides of the latest lecture will be on the course page. (<https://www.mit.bme.hu/eng/oktatas/targyak/vimiab00>)
These slides are under copyright.

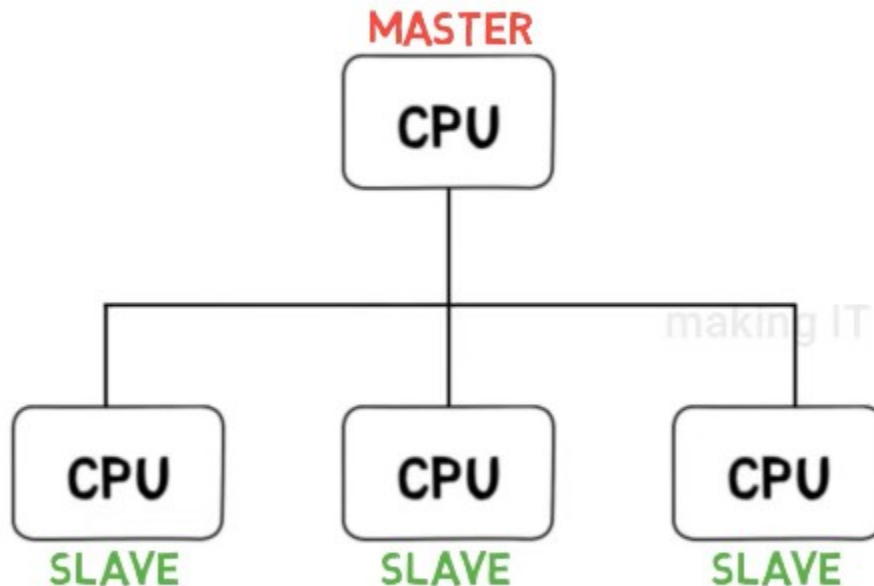
The main blocks of the OS and the kernel (recap)



Task scheduling (recap)

- Single level scheduling
 - FCFS, RR, SJF, SRTF, Priorities
 - Convey, Starvation
 - Measures: Avg. waiting time, turnaround time
- Multilevel scheduling
 - Multilevel static scheduling: fixed priorities (no queue change)
 - Multilevel dynamic scheduling: dynamic priorities (better in practice, because the tasks may change their nature)
- Multiprocessor scheduling
 - Processor affinity
 - Symmetric
 - Local queues for every processor
 - Asymmetric
 - Kernel gets a whole CPU

ASYMMETRIC




- 1) ONE IS CONSIDERED AS MASTER AND OTHERS ARE SLAVES
- 2) MASTER CPU DECIDES WHICH CPU WILL PERFORM WHICH TASK OR MASTER EXECUTES SYSTEM PROGRAMS AND SLAVES ARE USED FOR APPLICATION PROGRAMS
- 3) ONE PROCESSOR INTERACTS WITH I/O DEVICES AND OTHERS ARE USED FOR INCREASING PROCESSING POWER

ANY ARCHITECTURE WHERE ALL CPUS DONT HAVE SIMILAR RIGHTS IS CALLED ASYMMETRIC MULTI PROCESSING

1) SYMMETRIC MULTIPROCESSING

2) ASYMMETRIC MULTIPROCESSING

SYMMETRIC



1) EQUAL RIGHTS FOR ALL PROCESSORS

2) ANY CPU CAN ACCESS OR EXECUTE ANY PROCESS.

3) EACH CPU HAS ACCESS TO MEMORY I/O DEVICES OR ANY OTHER HARDWARE

ONE OPERATING SYSTEM CONTROLS ALL CPU AND ALL CPU HAVE EQUAL RIGHTS

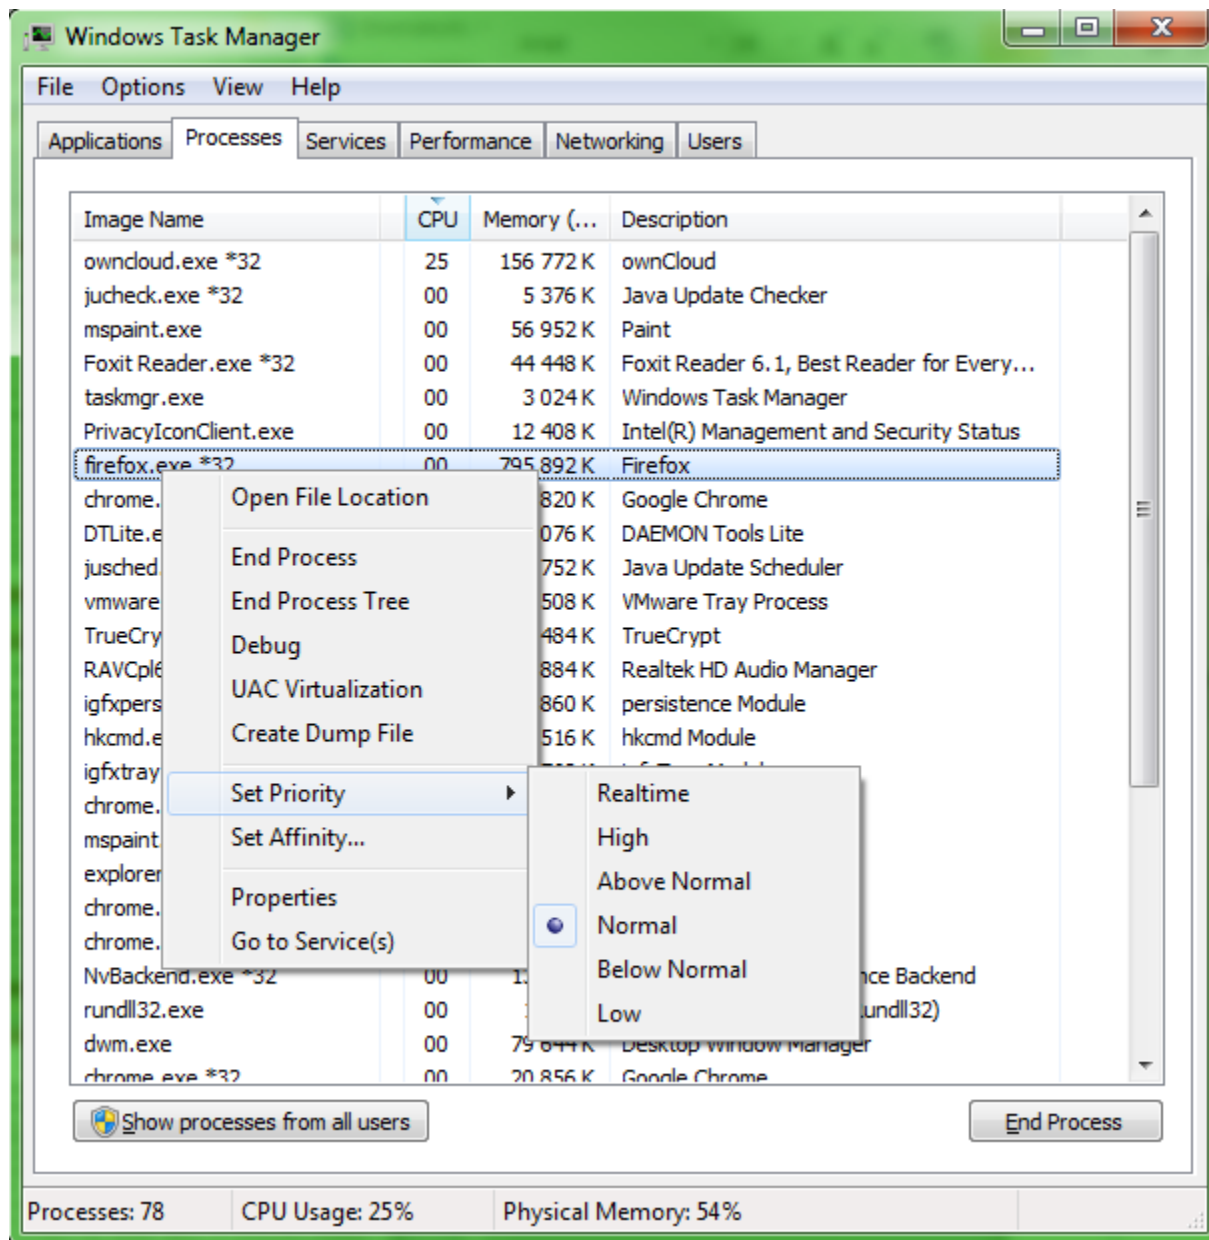
Advantages of Multi-processing:

Increased: 1- Throughput, 2- Reliability

3- Cost saving, 4- Parallel Processing

Schedulers in practice: Windows

- Symmetrical, multiprocessor with local RTR queues
 - It manages the soft and hard affinity
- Multilevel, priority, time sharing, fully preemptive [scheduler](#)
 - 16 „real-time” levels (priorities: 16-31)
 - 16 dynamic level (0-15, 0 is reserved for system)
 - The user can slightly modify the priorities
 - Low – below normal – normal – above normal – high – realtime
 - The scheduler may boost the priority in some cases
 - If a task finishes waiting for an I/O operation (better response time)
 - If a tasks gets an input from the user interface (better response time)
 - If a task is waiting for a long time (to avoid starvation)
 - Priority inversion is managed (see random boost)(Locks)
 - Multimedia Class Scheduler Service boosts the MM and gaming processes



Schedulers in practice: Linux

- Before v2: traditional UNIX scheduler (as seen before)
- Before v2.4 kernel
 - Scheduling classes: real-time, non-preemptive, normal
 - Complex scheduling with $O(N)$ (linear) complexity
 - Single scheduling queue (no multiprocessor support)
 - Kernel is non-preemptive
- Kernel v2.6
 - $O(1)$ scheduler (as the call it)
 - Distributed queues for each processor
 - Multilevel feedback with priorities and time-sharing
 - 140 priority level: 0-99 „real-time” static, 100-139 time-sharing, dynamic
 - „active” has time slice, „expired” (preempted) queues on every level
 - Priority is recalculated when moving from active to expired
 - If the active queue is empty, the expired become active
 - Tasks waiting for a long time gets a little „bonus”
- From kernel v2.6.23 CFS (Completely fair scheduler)
 - See next slide...

$O(1)$ --- constant complexity - not depending on

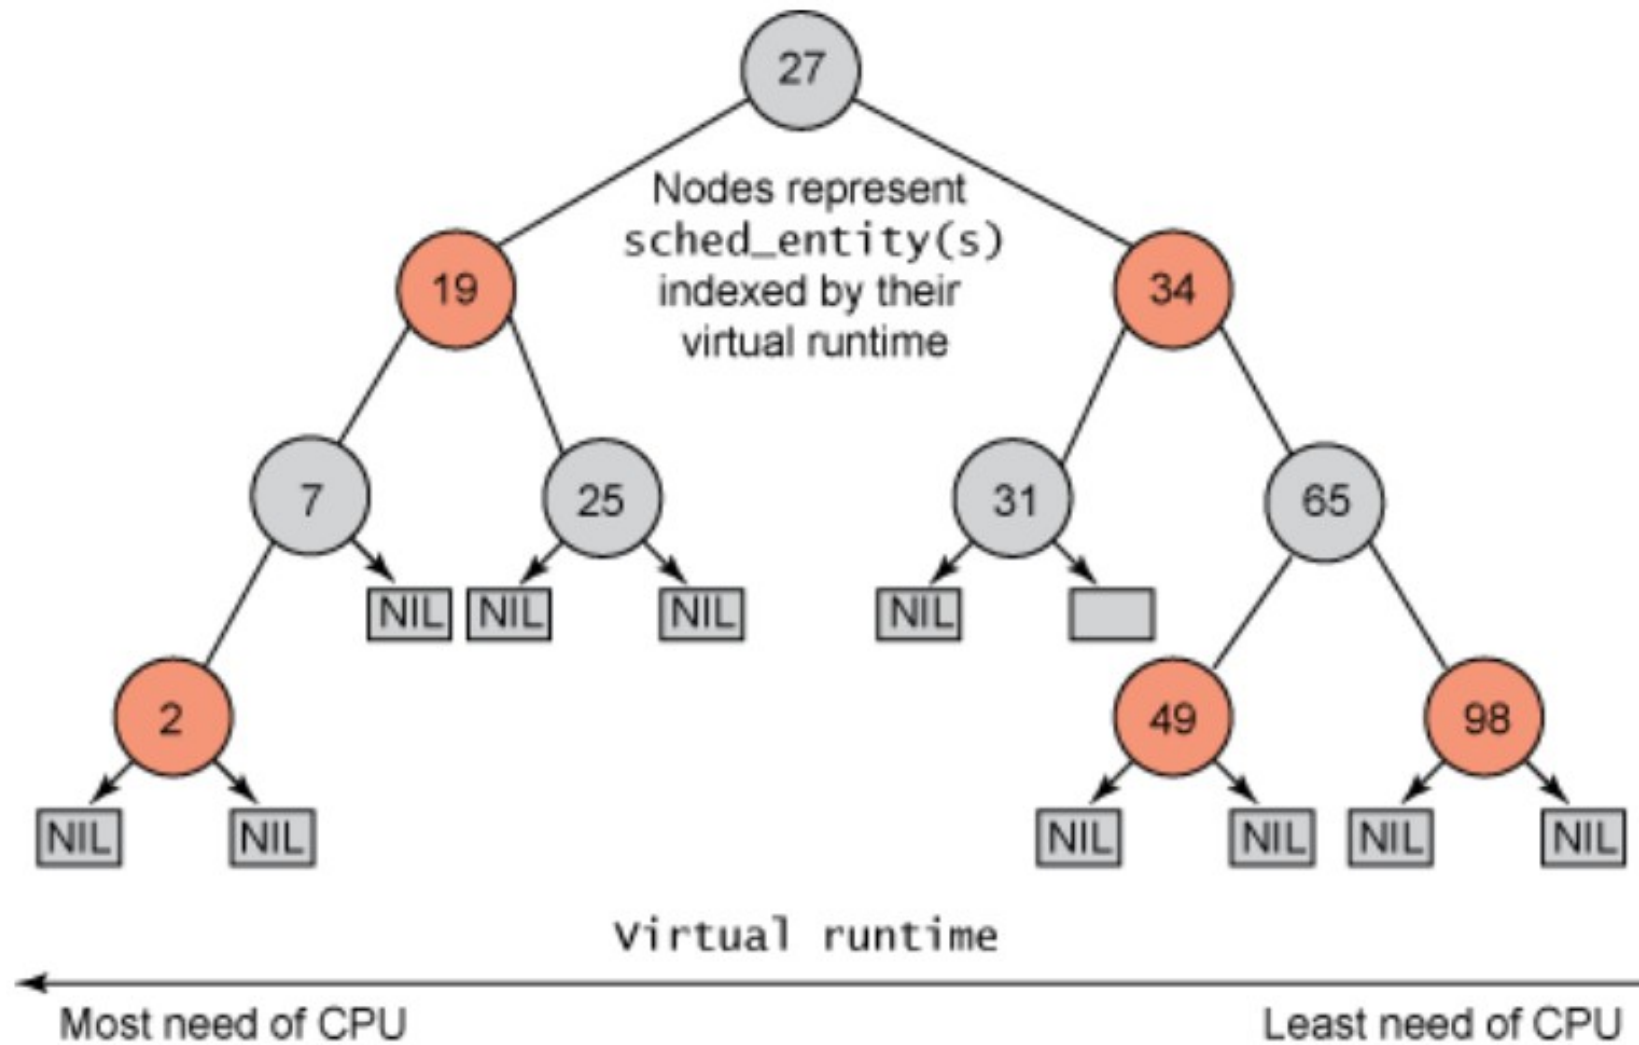
n

Linear $O(n)$ — performance becomes less

efficient as data grows.

Linux [CFS](#) (Completely Fair Scheduler)- Ingo Molnár

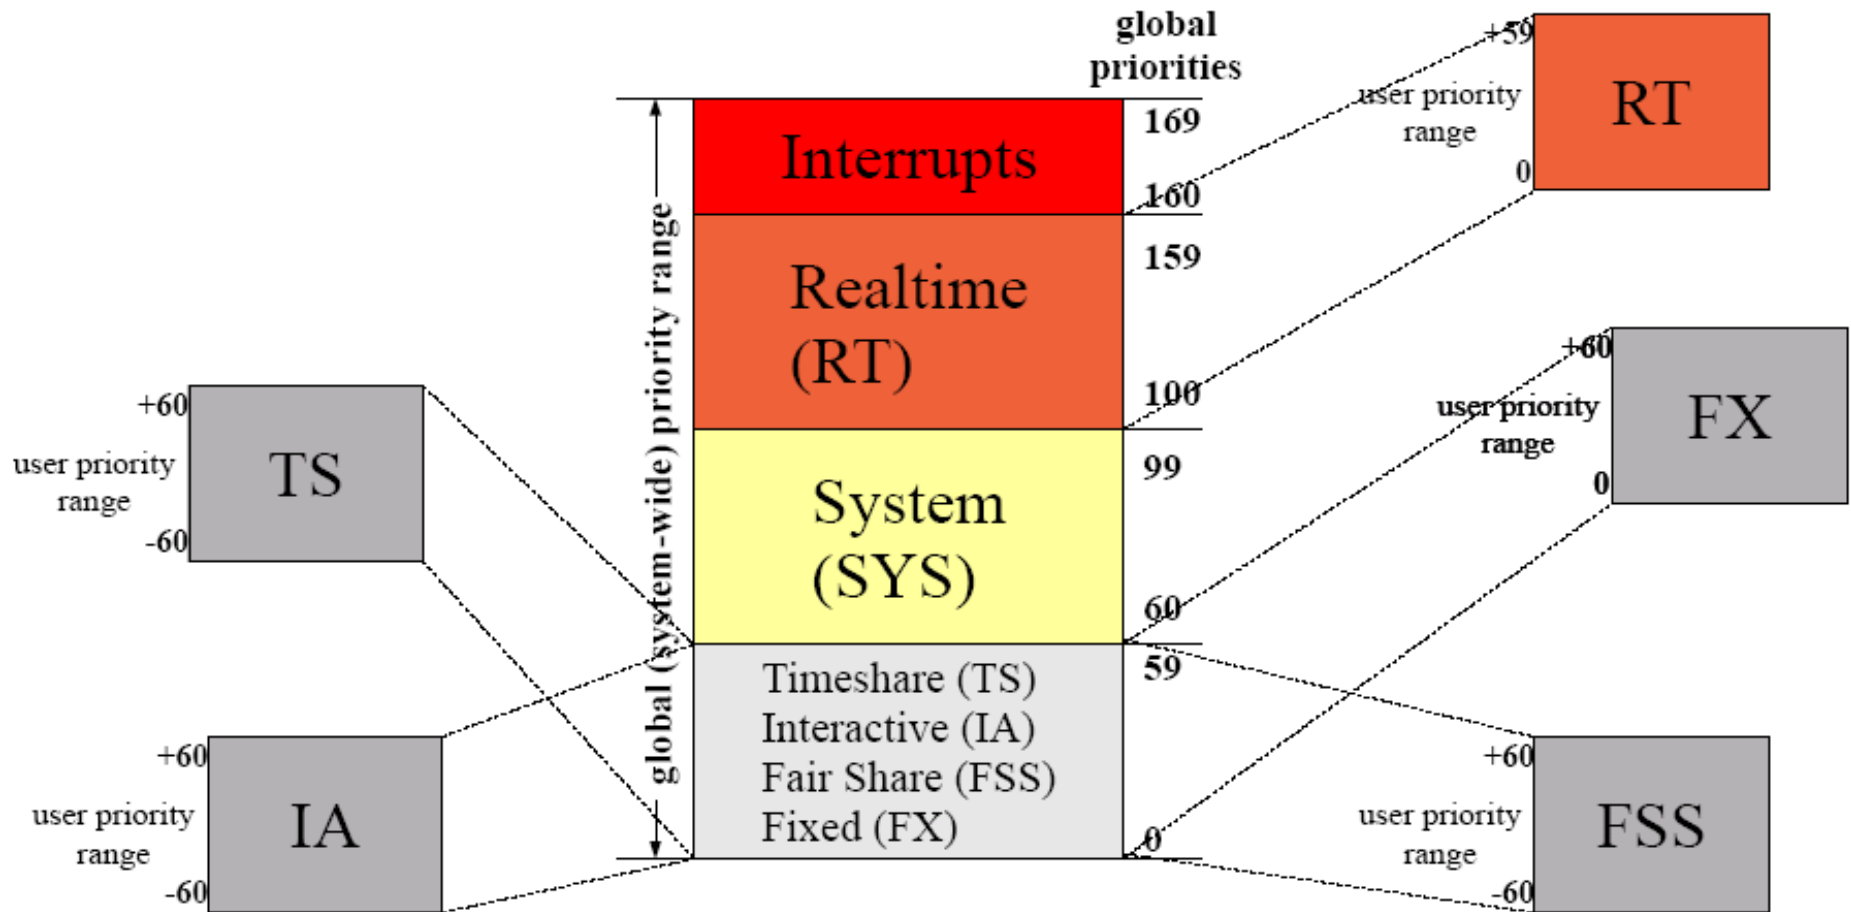
- Instead of queues, the tasks are stored in a special tree structure
 - Red-black tree: self balancing search tree
 - Smaller values on the left, higher values on the right
 - $O(\log n)$ Complexity
- Priority calculation
 - Virtual runtime (vruntime) assigned to each task
 - It tries to distribute runtime in a fair way
 - During execution vruntime is increasing, depending of the **priority**
 - The red-black tree is constructed based on vruntime
 - The tasks with
 - Small vruntime: gets the CPU earlier
 - if process has run for **t ms**, then
 $\text{vruntime} += t * (\text{weight based on nice of process})$



Solaris UNIX

- Properties
 - The scheduling is thread based
 - Fully preemptive kernel
 - Multiprocessor systems and virtualization are supported
- Multiple scheduling classes
 - Time-sharing (TS): scheduling by waiting/running times
 - Interactive (IA): like above, but enhance the priority of the active window
 - Fixed priority (FX)
 - Fair share (FSS): CPU resources are assigned to a process group
 - Real-time (RT): this level provides the shortest latency
 - Kernel threads (SYS): highest level except the RT

Scheduling classes of Solaris UNIX



Computation examples for simple schedulers

- Schedulers
 - FCFS: simple, based on FIFO (cooperative)
 - RR: time-sharing (after the time slice is up, gets the next task from FIFO)
 - SJF: ordering tasks by estimated CPU-burst (cooperative)
 - SRTF: preemptive SJF
 - PRI: ordering tasks by priority
- Measuring numbers for evaluation
 - Response time
 - Waiting time – elapsed time in **non** running state (waiting, ready to run)
 - Execution time – elapsed time in running state

Simple scheduler examples

- Show the operation of simple schedulers on Gantt chart and calculate the avg. waiting time!
 - Methods: FCFS, SJF, RR (TS=2, TS=6)
 - Tasks [Start, CPU-burst]
 - A [0, 6]
 - B [0, 3]
 - C [0, 3]

FCFS example

SJF example

RR (TS=2) example

RR (TS=6) example

Practice example for static multilevel scheduler

- The scheduler
 - Static priorities: 0..9 (0 is the highest)
 - Levels:
 - Level 1: 0-4 non preemptive SJF
 - Level 2: 5-9 preemptive RR, time slice: 2
 - To calculate: running order and avg. waiting time
- Parameters of the tasks
 - Priority
 - Starting (arrival) time
 - CPU-burst
- Tasks [Pri, Sta, CPU]
 - A [3, 0, 5]
 - B [4, 0, 4]
 - C [6, 5, 8]
 - D [6, 7, 8]

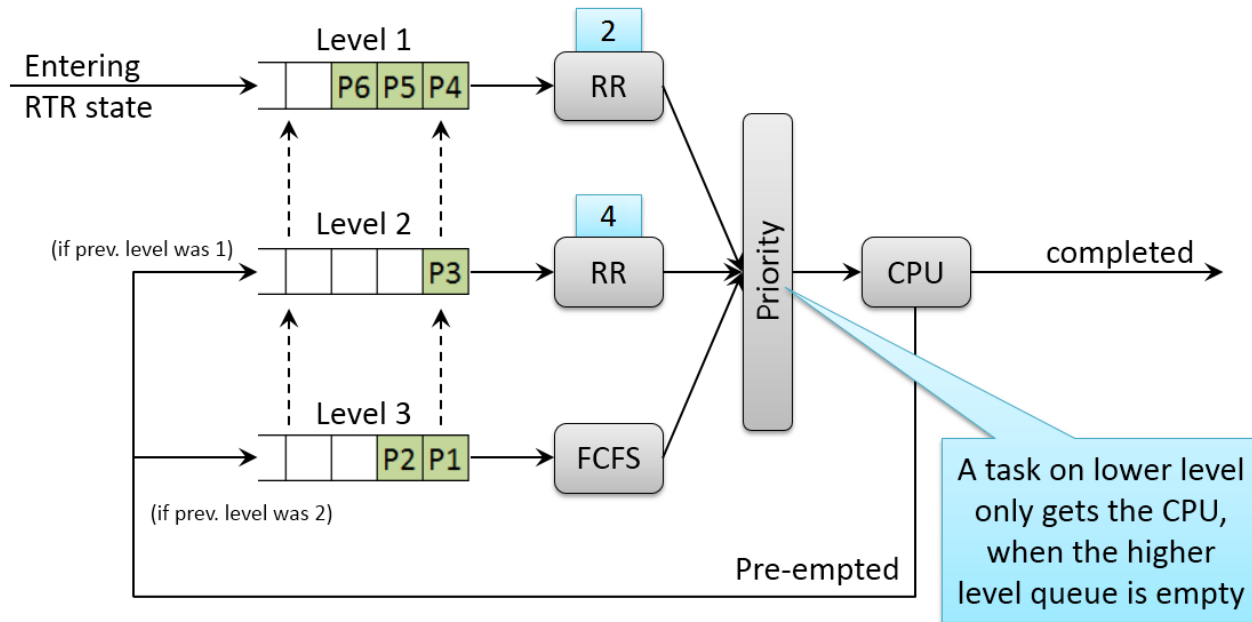
Gantt chart (static multilevel)

Tasks											
Task A [3, 0, 5]	A					A					
Task B [4, 0, 4]	B	B									
Task C [6, 5, 8]						C					
Task D [6, 7, 8]							D				
	Time steps	0	1	2	3	4	5	6	7	8	9

Tasks	The Gantt chart of the scheduling																										
Task A [3, 0, 5]	A					A																					
Task B [4, 0, 4]	B	B																									
Task C [6, 5, 8]						C				C			C			C			C			C					
Task D [6, 7, 8]						D				D			D			D			D			D			D		
Time steps	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26

MFQ (Multilevel Feedback Queue) example

- 3 levels:



- Tasks [Start, CPU-burst]

- A [0, 5]
- B [0, 5]
- C [3, 13]
- D [10, 11]

Gantt chart (MFQ)