

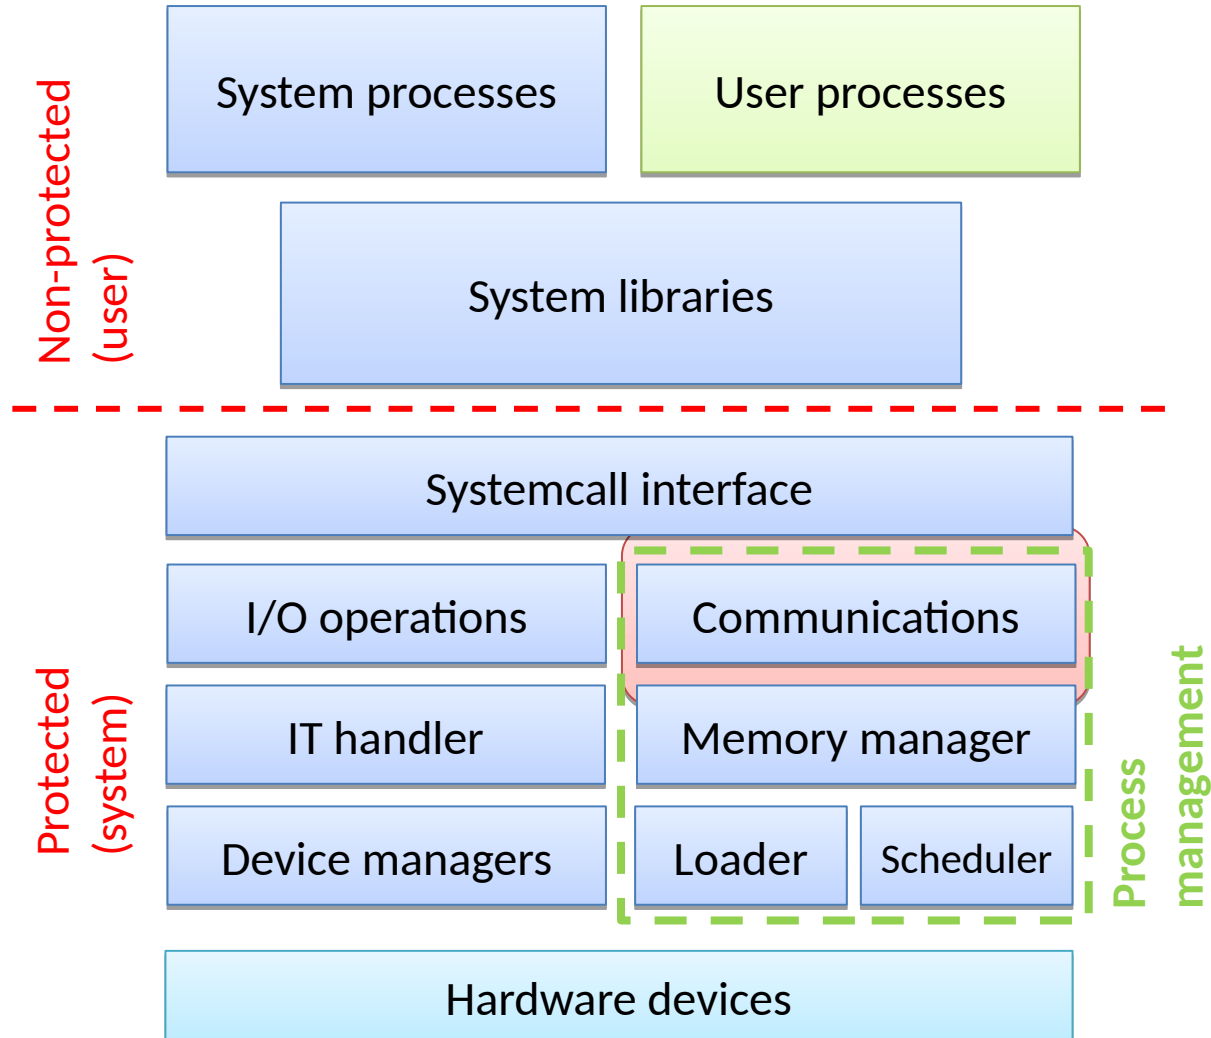
Operating Systems – Interprocess Communication

Hussein Al-Rikabi
rhussen@mit.bme.hu

Budapest University of Technology and Economics (BME)
Department of Measurement and Information Systems (MIT)

The slides of the latest lecture will be on the course page. (<https://www.mit.bme.hu/eng/oktatas/targyak/vimiab00>)
These slides are under copyright.

The main blocks of the OS and the kernel (recap)



The abstract virtual machine concept


- The main goal of the OS to support the execution of user jobs
 - Jobs are implemented by one or multiple tasks
 - Decomposed jobs may require communication between tasks
- Implementing tasks: processes and threads
 - Threads in the same process have a shared memory
 - Simple communication, but conflicts may arise
 - Processes are separated from each other
 - More complex communication
- Memory management separates the processes
 - Every process has its own memory range, which others cannot access
 - For efficient operation there can be overlapped ranges
 - Read-only pages can be used by multiple processes
 - The **Copy-On-Write** method accelerates the process creation

The basic forms of communication

- Through shared memory
 - The tasks memory range has a special range, which is translated to a shared physical memory
 - This shared range can be read and written by multiple tasks simultaneously
 - This method called: **PRAM model** (Pipelined RAM)
 - Competitive situations (e.g. multiple writers at the same time)
 - The execution order of these operations is not defined
- With messaging
 - No shared memory
 - The OS provides the communication mechanisms
 - Many solutions for this problem
 - Operations
 - `Send(recipient, data_pointer[, data_size]);`
 - `Receive(sender, data_pointer[, data_size]);`
 - Examples
 - Network communication, distributed systems, microkernel

The shared memory model (PRAM)


- Parallel tasks can use the shared memory range
 - They can start operations independently from each other
 - The operations may start at the same time
- Rules for the operations arrived at the same time
 - *Read-read*
 - Two simultaneous reads operations will both return the RAM contents
 - *Read-write*
 - The RAM will be written, the read op. will return the new value
 - *Write-write*
 - One of the two values will be written to the RAM



(b)

The shared memory model (PRAM)

- The effect of the rules
 - The operations has no effect to each other (no mix-up), hence the **pipelined** term: the operations are ordered in a pipeline
 - It's undetermined which op. will execute first if they arrive at the same time
 - To avoid undefined operations the tasks have to **synchronize** their operation



(b)

Communication methods based on the PRAM model

- Data exchange between the **threads** of the same process through global variables
 - Not supervised by the OS, the programmers design the operation
- Shared memory between **processes** (SHM)
 - The virtual memory range of the processes has a special part, which is translated to the same physical address
 - After the assignment, the usual memory operations can be used (no system calls, no additional overhead)
 - Fast method to communicate
 - The capacity of the buffer is limited (kernel defines the size)
 - Standard: POSIX Shared Memory
 - Implementations: System V UNIX, Windows (part of the kernel)
 - Typical applications: database engines

Producer-Consumer Problem

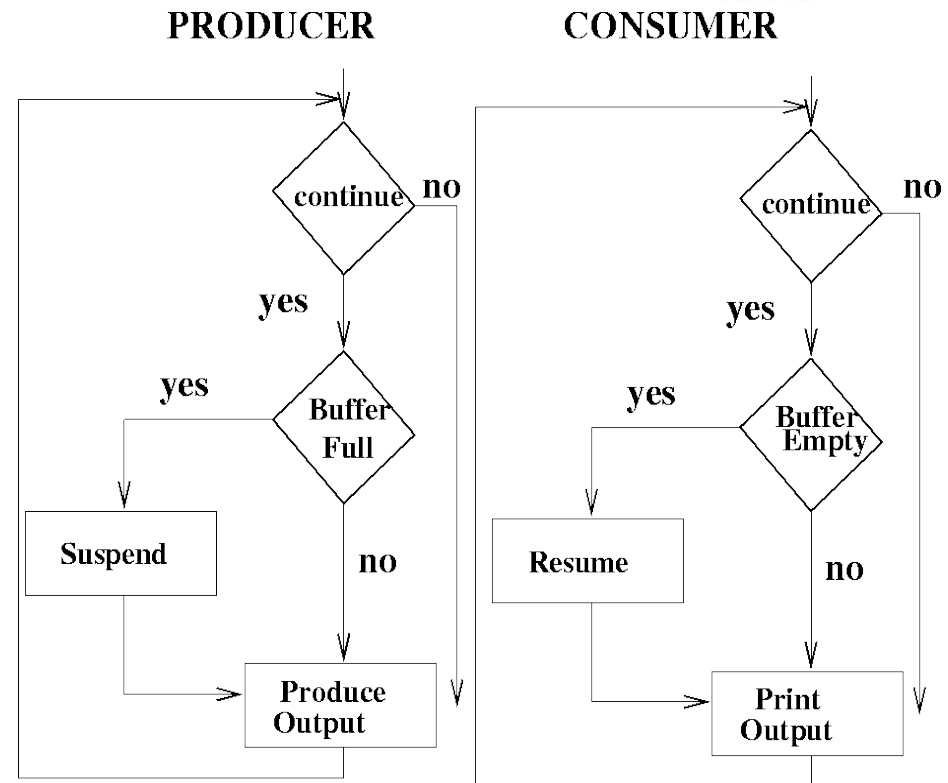

- Producer process produces information that is consumed by a consumer process.

- A **Buffer** is needed which can be:

- unbounded buffer
- bounded buffer

- What if Producer and Consumer tries to access the buffer simultaneously?

Synchronization



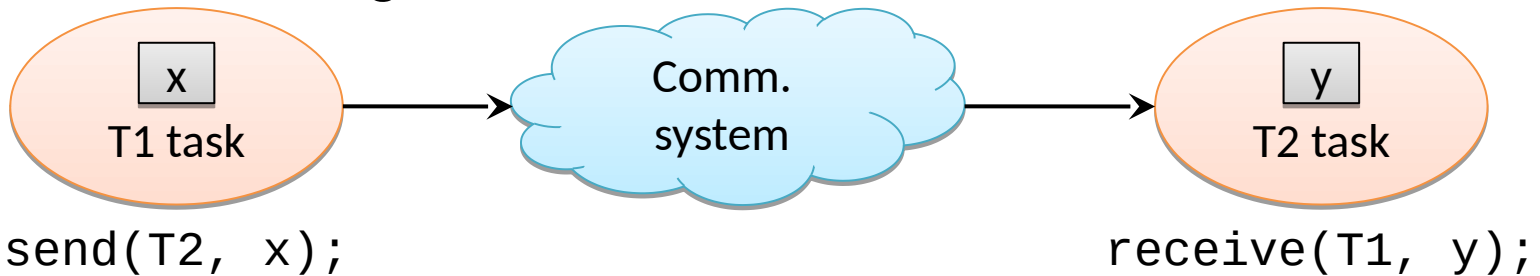
Communication with messaging

Basic questions of messaging (message-passing)

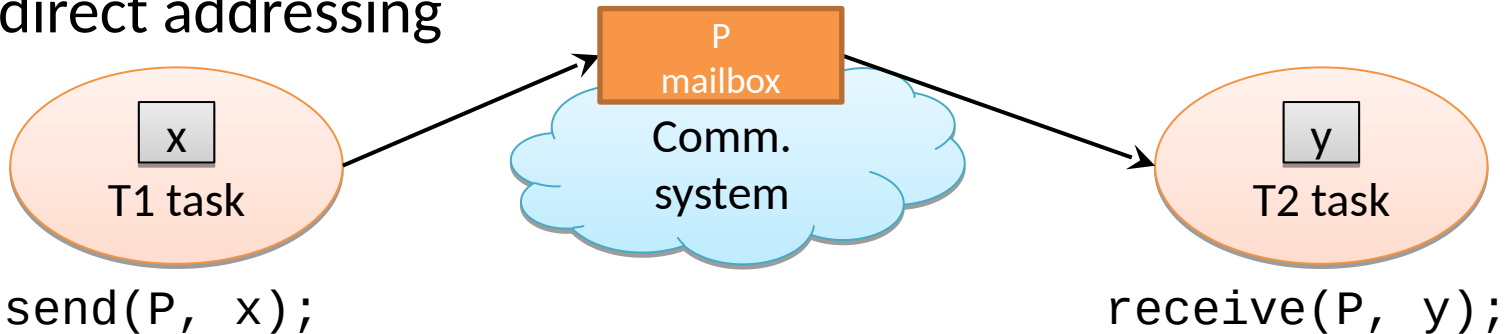
- Addressing: how can the recipient address determined?
 - Direct approach or using a mediator?
 - Single addressing or broadcasting?
 - Simultaneous receiving from multiple senders?
 - If the recipient knows the sender, can it decide whether to receive the message or not?
- Synchronization
 - Is the `send()` function blocking or non-blocking?
 - If not, how can the sender know that the message is received?
 - How the `receive()` function works?
 - What happens if the `receive()` function is called, but the data is not yet arrived?
 - Are acknowledgements mandatory, when the data is received?
- Semantics of data exchange
 - The data stored by the sender, or the receiver, or both?
- Performance, reliability
 - What throughput can be achieved? How much is the delay between send and receive?
 - How many and what size messages can be sent?
 - Who supervise (error handling) the process? What happens if an error occurs?

The basic methods of addressing


- Direct addressing



- Indirect addressing



- Asymmetric (direct on the sender's side) addressing



- Multiple addressing (multicast, broadcast)

Synchrony

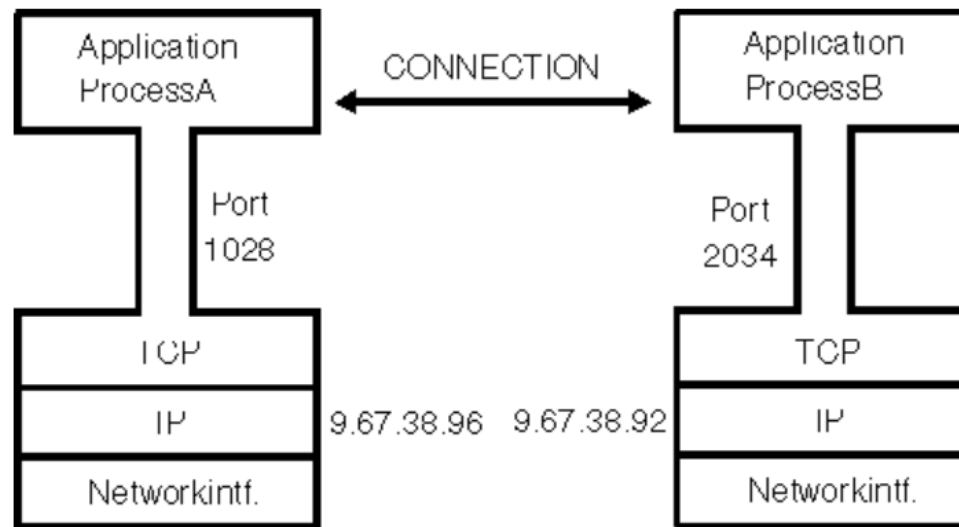
- **Synchronous data transfer (blocking)**
 - The operation blocks the task execution (waiting state)
 - Send: it returns only, when the data is received
 - If there are no mediator too long blocking is possible
 - Receive: it returns only, when the data is received
- **Asynchronous data transfer (non-blocking)**
 - The tasks are not blocked during the transfer
 - The results are not available when the function returns
 - The tasks has to be notified by other methods about the completion or the errors
 - The messages which are not received has to be stored
 - operations
 - Send: returns immediately after calling
 - The successful delivery has to be checked later
 - Receive: returns immediately with the data or with „no data received”
 - Busy waiting for a message should be avoided

Communication in Client–Server Systems

- Socket communication
 - Asymmetric method based on the client-server model
 - It is widely used on TCP/IP bases
 - The machines identified by IP addresses, the senders and receivers identified by ports
 - The OS provides the network connectors for the communication
- Remote Procedure Call (RPC)
 - Asymmetric method based on the client-server model
 - Calling a function in another task
 - The sender transfers the function name and arguments and receives the return value
 - The recipient performs the operations with the provided arguments and returns the results
 - It is based on network communications, so it can be used between different machines
 - Besides the communication protocol, the data semantics also determined
 - Exact data types, structures
 - Automatic conversion may be performed between the participants
 - Higher level OS services based on RPC
 - A number of development schemes (design patterns) are based on RPC

Network socket communications


- The communication is based on network protocols and addressing
 - The socket defines the protocol being used: TCP, UDP, or IP.
 - The IP and the port.
 - The identifier of the communication endpoint in tasks



Socket A = {TCP, 9.67.38.96, 1028}

Socket B = {TCP, 9.67.38.92, 2034}

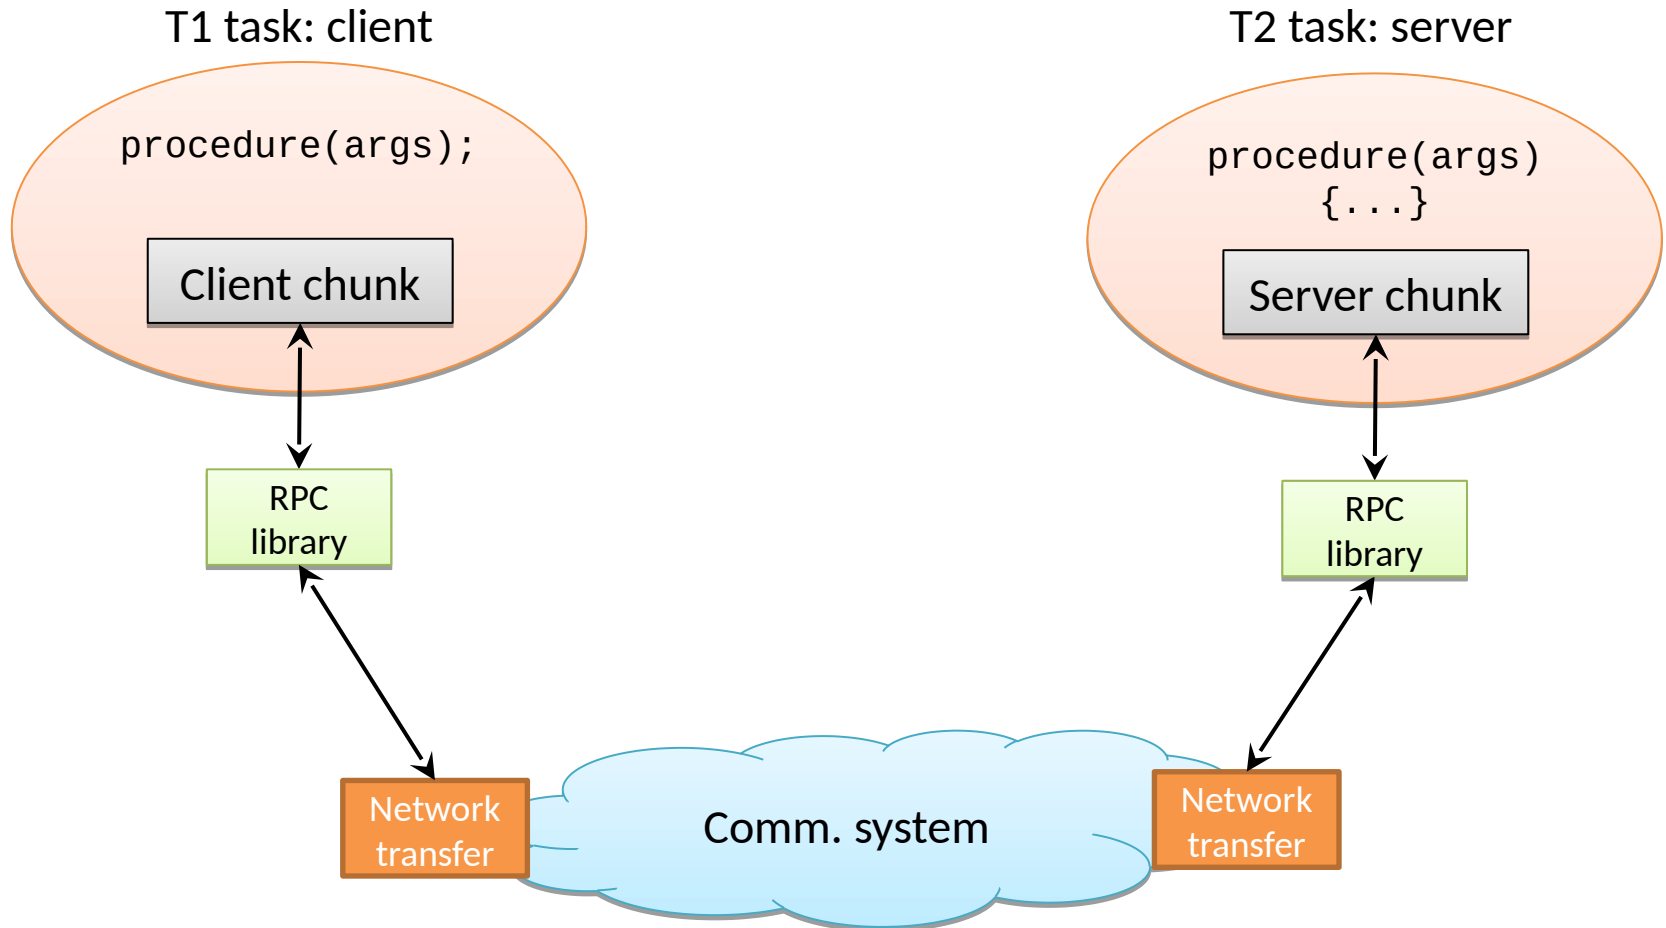
Network socket communications



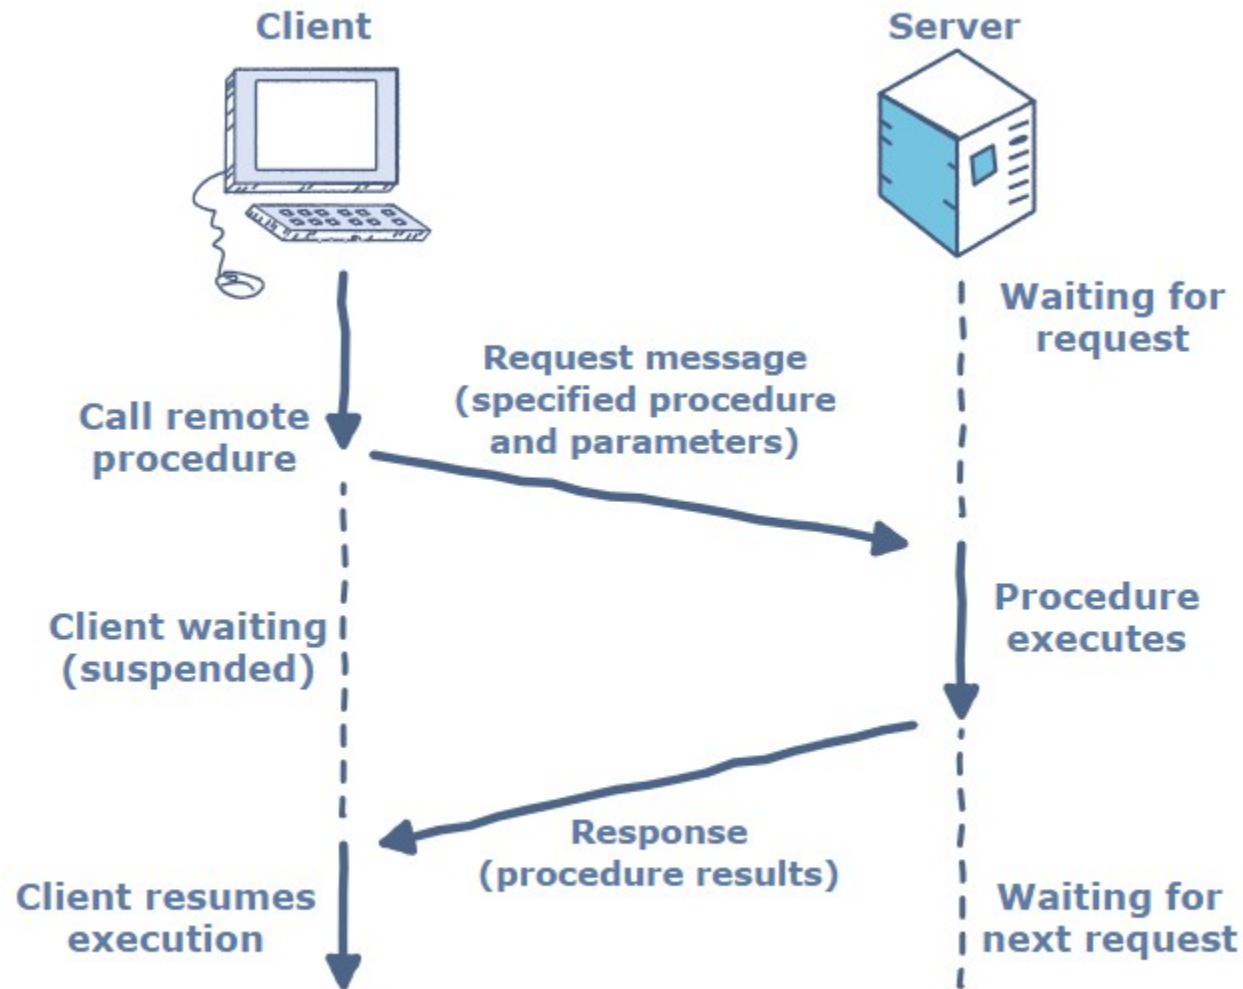
Remote Procedure Call (RPC)

- Distributed system infrastructure based on network communications
 - High level communication between processes
 - Remote function calls in different processes (even on a different machine)
 - Simple implementation due to common API-s
- Structures
 - Communication infrastructures
 - Protocols for transferring functions calls
 - Portmapper: assigning process ID-s and ports

The blocks of the RPC communication



RPC communication



A remote procedure call

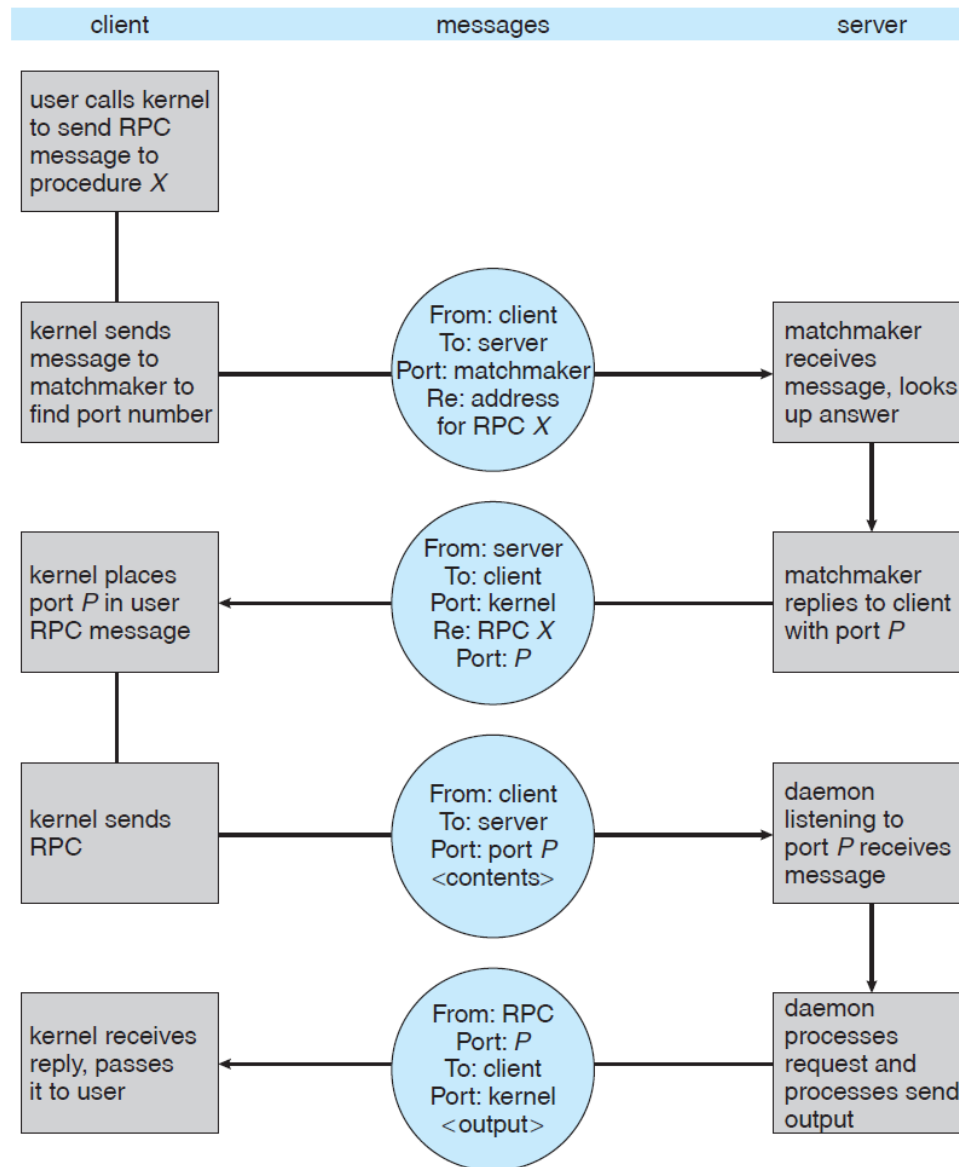


Figure 3.21 Execution of a remote procedure call (RPC).

The basic forms of communication (summary)

- **PRAM model**
 - Shared memory
 - The operations are pipelined
- Synchronous transfer operations
- No addressing
- Applications:
 - Threads between a process
 - **Shared memory between processes**
- Pro-s
 - Fast, simple
 - No overhead after initialization
- Con-s
 - R-W and W-W conflicts
 - Synchronization between participants is necessary
 - Limited capacity
- **Messaging systems**
 - It uses a communication system
 - Send and receive operations
 - Parallel operations are ordered
- Synchronous or asynchronous operation
- Addressing: direct, indirect, multi
- Applications
 - Mailbox
 - Pipeline
 - Message queue
 - **Network socket**
 - **Remote Procedure Call**
- Pro-s
 - Widely available
 - Between different machines
- Con-s
 - Communication errors has to managed
 - Slower

Interprocess communication in practice

POSIX Shared Memory Example

- Several IPC mechanisms are available for POSIX systems, including **shared memory** and message passing.
- A process creates a shared-memory --- `shmget()`
 - segment id = **shmget**(identifier, size, S_IRUSR | S_IWUSR);
 - identifier(key) - if set to IPC_PRIVATE, new shared-memory segment is created
 - size - the third parameter defines the **Mode** , read or write or both
- A successful call to **shmget()** returns an integer identifier, any task wants to access this memory, must specify this identifier.
- So a process can access it using **shmat(shmid, shmaddr , shmflg)**
 - shmaddr = the second parameter Points to the desired address of the shared memory segment
 - shmflg = Specifies a set of flags that indicate the specific shared memory conditions and options to implement.
- Also the share memory can be detached and removed, and there are many other commands to utilize it. See the code :

```
#include <stdio.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the identifier for the shared memory segment */
    int segment_id;
    /* a pointer to the shared memory segment */
    char *shared_memory;
    /* the size (in bytes) of the shared memory segment */
    const int size = 4096;

    /* allocate a shared memory segment */
    segment_id = shmget(IPC_PRIVATE, size, S_IRUSR | S_IWUSR);

    /* attach the shared memory segment */
    shared_memory = (char *) shmat(segment_id, NULL, 0);

    /* write a message to the shared memory segment */
    sprintf(shared_memory, "Hi there!");

    /* now print out the string from shared memory */
    printf("%s\n", shared_memory);

    /* now detach the shared memory segment */
    shmdt(shared_memory);

    /* now remove the shared memory segment */
    shmctl(segment_id, IPC_RMID, NULL);

    return 0;
}
```

Network Sockets message passing Example:


```
import java.net.*;
import java.io.*;
```

Importing required libraries

```
public class DateServer
{
```

Class : **server** or **clint** ?

```
    public static void main(String[] args) {
        try {
```

```
            ServerSocket sock = new ServerSocket(6013);
```

Creating Socket with Port No.

```
            // now listen for connections
```

```
            while (true) {
```

```
                Socket client = sock.accept();
```

Listening, waiting a clint

```
                PrintWriter pout = new
```

```
                    PrintWriter(client.getOutputStream(), true);
```

```
                // write the Date to the socket
```

```
                pout.println(new java.util.Date().toString());
```

Writing data to the socket

```
                // close the socket and resume
```

```
                // listening for connections
```

```
                client.close();
```

Close socket and wait again

```
            }
```

```
        }
```

```
        catch (IOException ioe) {
```

```
            System.err.println(ioe);
```

```
        }
```

```
    }
```

```
}
```

Some lines can be added here
to handle errors and exceptions

```

import java.net.*;
import java.io.*;

public class DateClient
{
    public static void main(String[] args) {
        try {
            //make connection to server socket
            Socket sock = new Socket("127.0.0.1",6013);

            InputStream in = sock.getInputStream();
            BufferedReader bin = new
                BufferedReader(new InputStreamReader(in));

            // read the date from the socket
            String line;
            while ( (line = bin.readLine()) != null)
                System.out.println(line);

            // close the socket connection
            sock.close();
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}

```

creates a Socket and requests a connection with the server

client define method of reading from socket, normal stream I/O statements

client can read from the socket

Socket is closed

End of Lecture